

Three experiment families answer three different questions

Study	Question and comparison	Measured endpoint	Interpretation
H1	Retention · W5, n=120 Which positions survive at k=0.25?	112/120 vs 1/120 fixed manifests and seeds	In-regime structural win
H2	INT4 quality · W3, n=240 Does role-aware placement beat uniform?	0.8833 vs 0.8792 fixed manifests and seeds	No quality separation
H3	Systems · H200, n=18 Mixed FlashInfer vs FullKV/SDPA	0.719× / 1.200–1.215× fixed manifests and seeds	Payload–latency trade-off